



A STEP WISE APPROACH TO CREATE A MACHINE LEARNING MODEL FOR AVOCADO DATASET


A Classification & Regression model building approach



Table of Contents

Introduction & Background	2
Problem Definition.....	3
Exploratory Data Analysis.	4
Initial Data Analysis.....	4
Descriptive Statistics of numerical data.....	6
Data Visualization	8
Pre-Processing Data.	17
Building Machine Learning Models.....	24

Introduction & Background



It is a well-known fact that Millennials LOVE Avocado Toast. It's also a well-known fact that all Millennials live in their parents' basements.

Clearly, they aren't buying home because they are buying too much Avocado Toast!

But maybe there's hope... if a Millennial could find a city with cheap avocados, they could live out the Millennial American Dream.

The aim of this blog is to help you understand how to create a Machine learning model & what are the necessary steps to predict the regions based on the dataset available at [this link](#).

(This blog is written with the assumption that the reader has preliminary understanding of machine learning terminologies & basic libraries machine learning libraries & intermediate level of understanding of python programming language.)

We will understand all the necessary steps needed to clean & manipulate data & build ML model. So, lets go & dive into it...

Problem Definition.

The dataset represents weekly 2018 retail scan data for National retail volume (units) and price. Retail scan data comes directly from retailers' cash registers based on actual retail sales of Hass avocados. Starting in 2013, the dataset reflects an expanded, multi-outlet retail data set. Multi-outlet reporting includes an aggregation of the following channels: grocery, mass, club, drug, dollar and military. The Average Price (of avocados) in the table reflects a per unit (per avocado) cost, even when multiple units (avocados) are sold in bags. The Product Lookup codes (PLU's) in the table are only for Hass avocados. Other varieties of avocados (e.g., green skins) are not included in this dataset.

Some relevant columns in the dataset:

- **Date** - The date of the observation
- **AveragePrice** - the average price of a single avocado
- **type** - conventional or organic
- **year** - the year
- **Region** - the city or region of the observation
- **Total Volume** - Total number of avocados sold
- **4046** - Total number of avocados with PLU 4046 sold
- **4225** - Total number of avocados with PLU 4225 sold
- **4770** - Total number of avocados with PLU 4770 sold

We will create one Classification model for predicting the region where avocados are sold & one regression model for predicting the average price for each avocado, type, year, volume, bag size and other features in the dataset.

We will import following machine learning libraries necessary for the model creation & validation;

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 %matplotlib inline
6 sns.set_style('darkgrid')
7
8 import statsmodels.api as sm
9 from collections import Counter
10
11 from sklearn.preprocessing import StandardScaler
12 from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold, GridSearchCV
13
14 #Logistic Regression
15 from sklearn.linear_model import LogisticRegression
16
17 # Random Forest Classifier & Gradient Boosting Classifier
18 from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier
19
20 # Decision Tree Classifier
21 from sklearn.tree import DecisionTreeClassifier
22
23 # K Neighbors Classifier
24 from sklearn.neighbors import KNeighborsClassifier
25
26 #Classification model Evaluation metrics
27 from sklearn.metrics import confusion_matrix, classification_report, f1_score
28 from sklearn.metrics import roc_auc_score, accuracy_score, roc_curve, auc
29
30 import pickle
31
32 import warnings
33 warnings.filterwarnings('ignore')
34
35 #regression models
36 from sklearn.linear_model import LinearRegression, Ridge, Lasso
37 from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor, GradientBoostingRegressor
38 from sklearn.tree import DecisionTreeRegressor
39 from sklearn.neighbors import KNeighborsRegressor
40 from xgboost import XGBRegressor
41
42
43 #Regression evaluation metrics
44 from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

Exploratory Data Analysis.

Initial Data Analysis

- We will load the data using pandas 'pd.read_csv()' method.

```
1 data_url = "avocado.csv"
2 data0 = pd.read_csv(data_url)
```

- Reading the first & last 5 rows of data

```
1 data0.head()
```

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags	XLarge Bags	type	year	region
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62	93.25	0.0	conventional	2015	Albany
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07	97.49	0.0	conventional	2015	Albany
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042.21	103.14	0.0	conventional	2015	Albany
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5677.40	133.76	0.0	conventional	2015	Albany
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986.26	197.69	0.0	conventional	2015	Albany

```
1 data0.tail()
```

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags	XLarge Bags	type	year	region
18244	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	13498.67	13066.82	431.85	0.0	organic	2018	WestTexNewMexico
18245	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	9264.84	8940.04	324.80	0.0	organic	2018	WestTexNewMexico
18246	9	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	9394.11	9351.80	42.31	0.0	organic	2018	WestTexNewMexico
18247	10	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	10969.54	10919.54	50.00	0.0	organic	2018	WestTexNewMexico
18248	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	12014.15	11988.14	26.01	0.0	organic	2018	WestTexNewMexico

Observation:

Here we see there are total 14 columns, out of which 'region' is our target/label column or dependent variable. We also see that there is a column named 'Unnamed:', it is an observation number from a specific region (you can check this using data0.head(55) & data0.sample(15)). Observations are taken on every 7th day from the last observation. We will remove this column as it brings no meaning our model.

Let's understand our label column,

```
1 data0.region.nunique()
```

54

There are total 54 unique entries in target column, following are the values;

```
1 data0.region.unique()
```

```
array(['Albany', 'Atlanta', 'BaltimoreWashington', 'Boise', 'Boston',  
      'BuffaloRochester', 'California', 'Charlotte', 'Chicago',  
      'CincinnatiDayton', 'Columbus', 'DallasFtWorth', 'Denver',  
      'Detroit', 'GrandRapids', 'GreatLakes', 'HarrisburgScranton',  
      'HartfordSpringfield', 'Houston', 'Indianapolis', 'Jacksonville',  
      'LasVegas', 'LosAngeles', 'Louisville', 'MiamiFtLauderdale',  
      'Midsouth', 'Nashville', 'NewOrleansMobile', 'NewYork',  
      'Northeast', 'NorthernNewEngland', 'Orlando', 'Philadelphia',  
      'PhoenixTucson', 'Pittsburgh', 'Plains', 'Portland',  
      'RaleighGreensboro', 'RichmondNorfolk', 'Roanoke', 'Sacramento',  
      'SanDiego', 'SanFrancisco', 'Seattle', 'SouthCarolina',  
      'SouthCentral', 'Southeast', 'Spokane', 'StLouis', 'Syracuse',  
      'Tampa', 'TotalUS', 'West', 'WestTexNewMexico'], dtype=object)
```

Observation:

From above we see that, amongst the city name we also see there is an entry as 'TotalUS'. this means some observations were entered using total of all the regions in US. These entries should be dropped as we are predicting individual regions.

Going ahead with the next step we will check missing entries, data types, total number of unique entries in all features, column names & data size;

```
1 data0.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 18249 entries, 0 to 18248  
Data columns (total 14 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   Unnamed: 0             18249 non-null   int64  
1   Date                   18249 non-null   object  
2   AveragePrice           18249 non-null   float64  
3   Total Volume           18249 non-null   float64  
4   4046                    18249 non-null   float64  
5   4225                    18249 non-null   float64  
6   4770                    18249 non-null   float64  
7   Total Bags              18249 non-null   float64  
8   Small Bags             18249 non-null   float64  
9   Large Bags             18249 non-null   float64  
10  XLarge Bags            18249 non-null   float64  
11  type                   18249 non-null   object  
12  year                   18249 non-null   int64  
13  region                 18249 non-null   object  
dtypes: float64(9), int64(2), object(3)  
memory usage: 1.9+ MB
```

```
1 data0.nunique()
```

```
Unnamed: 0      53  
Date            169  
AveragePrice    259  
Total Volume    18237  
4046            17702  
4225            18103  
4770            12071  
Total Bags      18097  
Small Bags      17321  
Large Bags      15082  
XLarge Bags     5588  
type            2  
year            4  
region          54  
dtype: int64
```

```
1 data0.columns
```

```
Index(['Unnamed: 0', 'Date', 'AveragePrice', 'Total Volume', '4046', '4225',  
      '4770', 'Total Bags', 'Small Bags', 'Large Bags', 'XLarge Bags', 'type',  
      'year', 'region'],  
      dtype='object')
```

```
1 data0.shape
```

```
(18249, 14)
```

Observation:

- There are total **18249 observations** in dataset & **14 columns**, & at initial level there are no missing entries.
- **Date** is formatted as object datatype, we will convert it into week number & month number, which will be helpful for visualization. The data was gathered during 169 unique days spanned over **4 years**.
- **Average price** is continuous type of numerical data.
- **type** is an object data type & it has 2 unique entries.
- type, year & region are the 3 categorical features from the dataset. from this we will remove year column as date is also there in the dataset.
- except Unnamed: 0, Date, type year & region all other features are float type continuous data.

Descriptive Statistics of numerical data

```
1 #Descriptive statistics
2 data0.drop(columns = ['Unnamed: 0', 'year', 'Date', 'type', 'region']).describe().T
```

	count	mean	std	min	25%	50%	75%	max
AveragePrice	18249.0	1.405978	4.026766e-01	0.44	1.10	1.37	1.66	3.25
Total Volume	18249.0	850644.013009	3.453545e+06	84.56	10838.58	107376.76	432962.29	62505646.52
4046	18249.0	293008.424531	1.264989e+06	0.00	854.07	8645.30	111020.20	22743616.17
4225	18249.0	295154.568356	1.204120e+06	0.00	3008.78	29061.02	150206.86	20470572.61
4770	18249.0	22839.735993	1.074641e+05	0.00	0.00	184.99	6243.42	2546439.11
Total Bags	18249.0	239639.202060	9.862424e+05	0.00	5088.64	39743.83	110783.37	19373134.37
Small Bags	18249.0	182194.686696	7.461785e+05	0.00	2849.42	26362.82	83337.67	13384586.80
Large Bags	18249.0	54338.088145	2.439660e+05	0.00	127.47	2647.71	22029.25	5719096.61
XLarge Bags	18249.0	3106.426507	1.769289e+04	0.00	0.00	0.00	132.50	551693.65

Observation:

- We see no missing observation from count. every feature has 18249 number of entries.
- There are some entries/observations where the minimum value of features was 0.0.

```
1 #percentage difference between mean & median
2 (abs((data0.drop(columns = ['Unnamed: 0', 'year', 'Date', 'type', 'region']).describe().T['mean'] -
3 data0.drop(columns = ['Unnamed: 0', 'year', 'Date', 'type', 'region']).describe().T['50%'])*
4 100/data0.drop(columns = ['Unnamed: 0', 'year', 'Date', 'type', 'region']).describe().T['mean'])).sort_values()
```

```
AveragePrice      2.558959
Total Bags         83.415138
Small Bags        85.530412
Total Volume      87.377004
4225              90.153966
Large Bags        95.127341
4046              97.049471
4770              99.190052
XLarge Bags      100.000000
dtype: float64
```

```
1 #percentage difference between 75% quantile & max
2 ((data0.drop(columns = ['Unnamed: 0', 'year', 'Date', 'type', 'region']).describe().T['max'] -
3 data0.drop(columns = ['Unnamed: 0', 'year', 'Date', 'type', 'region']).describe().T['75%'])*
4 100/data0.drop(columns = ['Unnamed: 0', 'year', 'Date', 'type', 'region']).describe().T['max'])).sort_values()
```

```
AveragePrice      48.923077
4225              99.266230
Total Volume      99.307323
Small Bags        99.377361
Total Bags        99.428160
4046              99.511862
Large Bags        99.614812
4770              99.754818
XLarge Bags       99.975983
dtype: float64
```

Observation:

- The mean & median % difference for all columns is more than 50% implying skewness in the dataset, except AveragePrice.
- features with % difference between 75% quantile & maximum value, more than 99% indicate possible outliers.

Dropping the 'Unnamed: 0' column & checking for duplicate & checking zero entry count.

Checking for duplicate entries...

```
1 #dropping the unnamed & year column from the dataset.
2 data0.drop('Unnamed: 0', axis = 1, inplace = True)
```

```
1 data0.duplicated().sum()
```

0

Observations:

1. No Duplicate entries.

Checking entries with ZERO Value...

```
1 data0[data0 == 0.0].count()*100/data0.shape[0]
```

```
Date                0
AveragePrice        0
Total Volume        0
4046                242
4225                61
4770                5497
Total Bags          15
Small Bags          159
Large Bags          2370
XLarge Bags        12048
type                0
year                0
region              0
dtype: int64
```

Observation:

1. PLU 4046, 4225, & 4770 has 242, 61 & 5497 entries as zero respectively.
2. XLarge bags has 12048 entries with zero value, highest of the all features. Small, large & total bags also have zero entries.

We will rename all the columns to be able to easily understand & remove whitespaces. (This step is not necessary but I like to do it, whenever I see whitespace in column names)

renaming column names to remove whitespaces...

```
1 columns = data0.columns.tolist()
2 columns_renamed = [column.strip().replace(" ", "_").lower() for column in columns]
3 data0.columns = columns_renamed
4 data0.rename(columns = {'4046':'plu_4046', '4225':'plu_4225', '4770':'plu_4770'}, inplace = True)
5 data0.columns
```

```
Index(['date', 'averageprice', 'total_volume', 'plu_4046', 'plu_4225',
      'plu_4770', 'total_bags', 'small_bags', 'large_bags', 'xlarge_bags',
      'type', 'year', 'region'],
      dtype='object')
```

Now we will copy the input data frame & do feature engineering for Date column on the copied dataframe as follows:


```

1 #Copying dataframe
2 data1 = data0.copy()
3
4 data1['date'] = pd.to_datetime(data1['date'], errors='raise')
5 data1['week_number'] = data1['date'].dt.week
6 data1['month_number'] = data1['date'].dt.month
7 data1['year'] = data1['date'].dt.year
8
9 data1.drop('date', axis=1, inplace=True)

```

The new dataframe looks like as follows:

1	data1													
	averageprice	total_volume	plu_4046	plu_4225	plu_4770	total_bags	small_bags	large_bags	xlarge_bags	type	year	region	week_number	month_number
0	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62	93.25	0.0	conventional	2015	Albany	52	12
1	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07	97.49	0.0	conventional	2015	Albany	51	12
2	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042.21	103.14	0.0	conventional	2015	Albany	50	12
3	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5677.40	133.76	0.0	conventional	2015	Albany	49	12
4	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986.26	197.69	0.0	conventional	2015	Albany	48	11
...
18244	1.63	17074.83	2046.96	1529.20	0.00	13498.67	13066.82	431.85	0.0	organic	2018	WestTexNewMexico	5	2
18245	1.71	13888.04	1191.70	3431.50	0.00	9264.84	8940.04	324.80	0.0	organic	2018	WestTexNewMexico	4	1
18246	1.87	13766.76	1191.92	2452.79	727.94	9394.11	9351.80	42.31	0.0	organic	2018	WestTexNewMexico	3	1
18247	1.93	16205.22	1527.63	2981.04	727.01	10969.54	10919.54	50.00	0.0	organic	2018	WestTexNewMexico	2	1
18248	1.62	17489.58	2894.77	2356.13	224.53	12014.15	11988.14	26.01	0.0	organic	2018	WestTexNewMexico	1	1

18249 rows x 14 columns

Data Visualization

We will make two group for data visualization purpose;

Numerical features :

1. averageprice
2. total_volume
3. plu_4046
4. plu_4225
5. plu_4770
6. total_bags
7. small_bags
8. large_bags
9. xlarge_bags

Categorical features :

1. type
2. year
3. region
4. week_number
5. month_number

Converting Some of the features in to proper data type for visualization purpose. & Creating two dataframes for numerical features & categorical features;

```

1 data1['year'] = data1['year'].astype('object')
2 data1['week_number'] = data1['week_number'].astype('object')
3 data1['month_number'] = data1['month_number'].astype('object')

1 month_dict = {1:'January',2:'February',3:'March',4:'April',5:'May',6:'June',
2               7:'July',8:'August',9:'September',10:'October',11:'November',12:'December'}

1 data1['month_number'] = data1['month_number'].map(month_dict)

1 data1_num = data1[num_cols]
2 data1_cat = data1[cat_cols]

```


Univariate Analysis

For Numerical features:

```

1 #distribution of numerical features..
2 i=0
3 plt.figure(figsize=(14,12))
4 for column in data1_num.columns:
5     plt.subplot(3,3,i+1)
6     sns.distplot(data1_num[column], bins=40, label="Skewness : %.2f"%(data1_num[column].skew()))
7     plt.xlabel(column, fontsize=12)
8     i+=1
9
10 plt.tight_layout()


```



Observation:

- All the numerical features except average price are skewed & it is evident from the above distribution plots.

For categorical features:

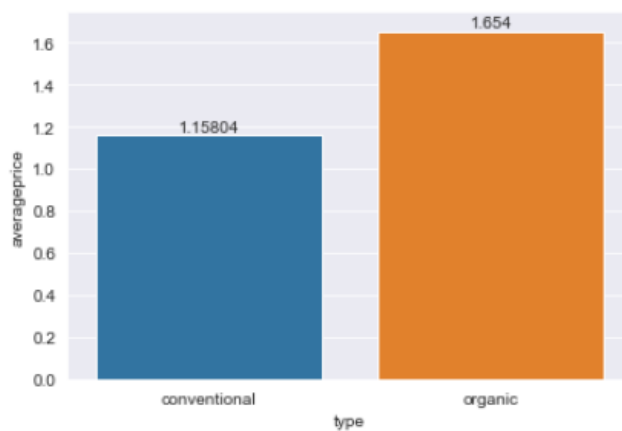


Observation:

- We can say that Conventional type Avocados are cheaper than the Organic type of Avocado.
- Organic type of avocado has average price ranging from 0.4 to 3.3, with mean value close to 1.65.
- Conventional type of avocado has avg price within 0.4 to 2.4, & mean being closer to 1.15.

Average price of avocado sold of each type

```
type
conventional    1.158040
organic         1.653999
Name: averageprice, dtype: float64
```

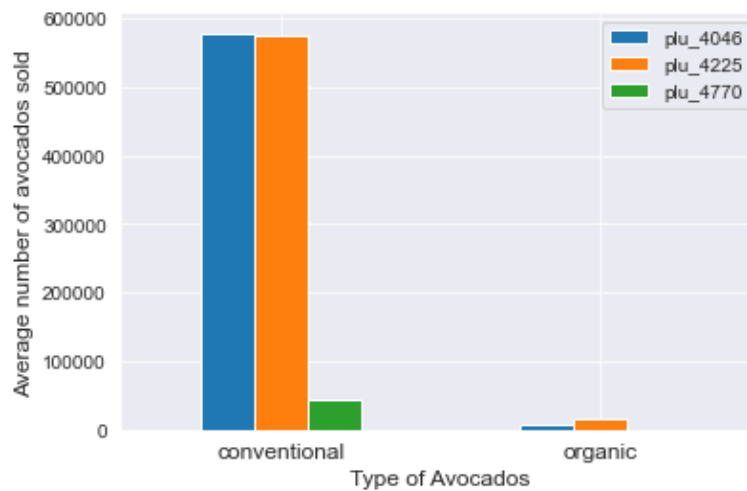


Observation:

- An organic type of Avocado is about 50 cents costlier than conventional type.

Bivariate Analysis


1. Average number of avocados sold based on product lookup codes of each type;



	plu_4046	plu_4225	plu_4770
type			
conventional	578611.649925	574805.318859	45405.796798
organic	7311.281600	15411.857724	266.254582

Observation:


- Mostly Conventional type of avocados with package label 4046 & 4225 are sold.
2. Average number of avocados sold based on bags size of each year



Observation:

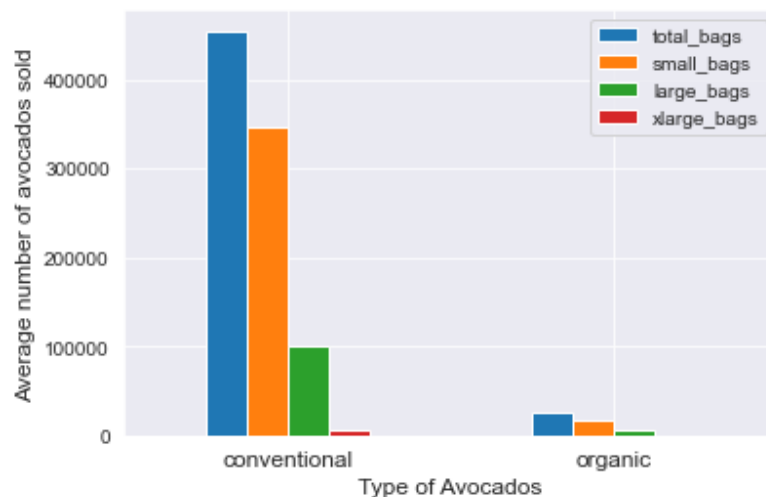
- In all year of the data available, we see Conventional type of avocados with package label 4046 & 4225 are sold most in terms of volumes.

3. Average volume of avocado sold in different regions based on package label type;



observation:

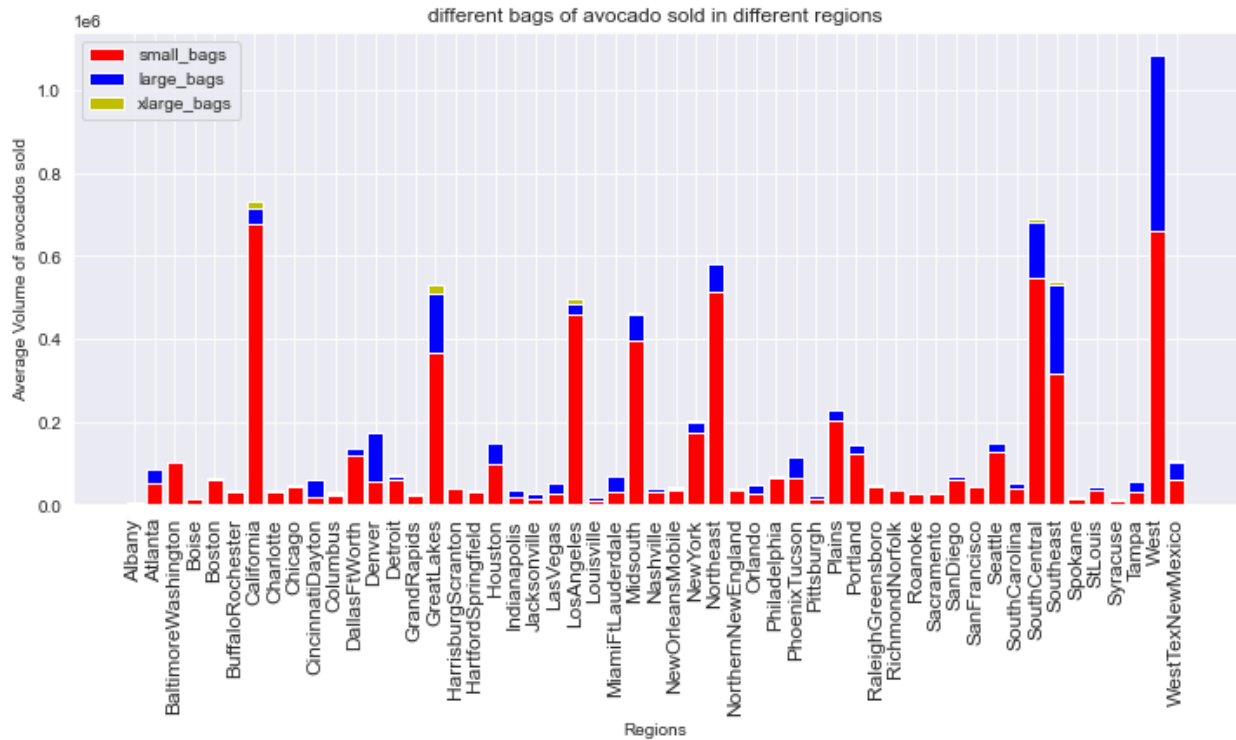
- California, South central, South east & west regions are the regions with most consumption of avocados. i.e., higher demand from these 4 regions.
4. Average number of avocados sold based on bags size of each type;



observation:

- Organic type of avocado is very less consumptive but higher price.
- Conventional type of avocado is sold heavily with higher volume & lower price.
- Smaller bags of conventional avocado is sold in much bigger volume than other bag sizes.


5. Different bags of avocado sold in different regions based on bag type;




Observation:

- Smaller bags of avocados are sold heavily in all the regions.


6. Average price of each type of avocado sold each year;



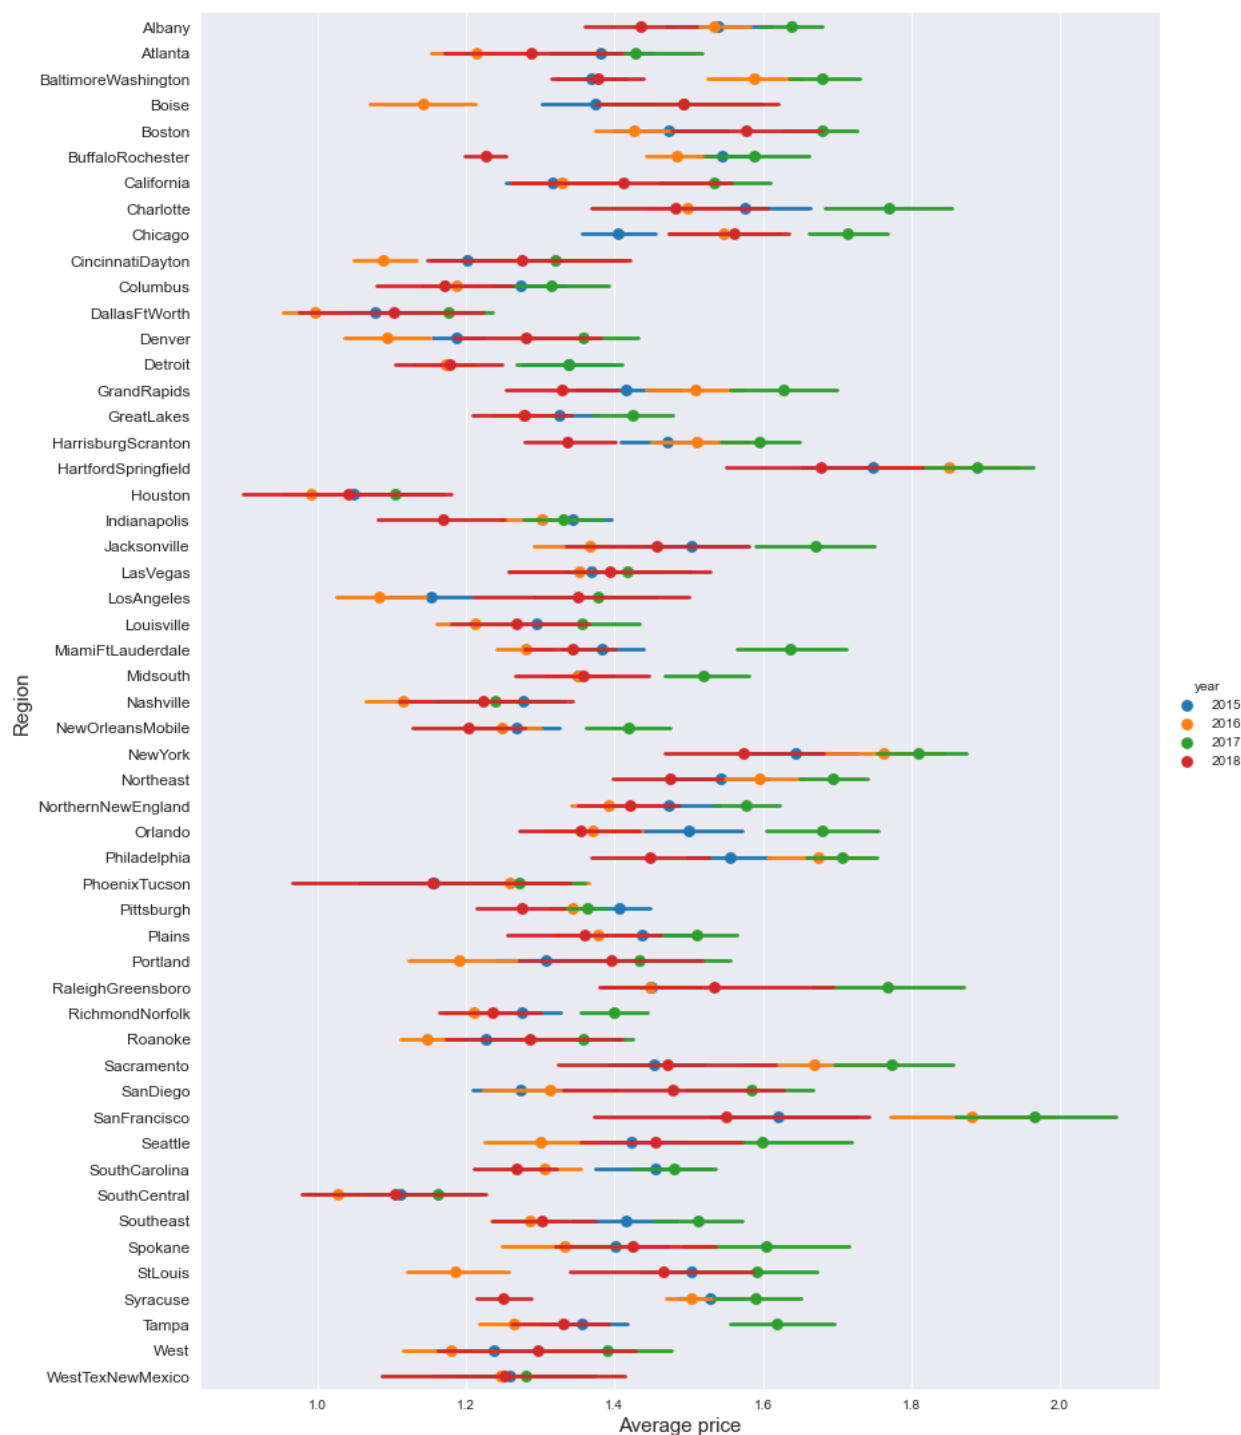
7. Average number of each type of avocado sold each year;



8. Average price variation during each year



9. Average price variation in all regions during each year



In 2015:

- average price of avocado was higher in Hart fort Springfield, New York, San Francisco, charlotte, philly. price was above 1.5 per unit.
- Houston, Dallas. South Central & LA had lower price per unit.
- Volume consumption was higher in California, West & South-central regions

In 2016:

- average price of avocado was higher in Hart fort Springfield, New York, San Francisco, philly, Sacramento & Northeast regions. price was above 1.6 per unit.
- Houston, Dallas. South Central & LA had lower price per unit.
- Volume consumption was higher in California, West & South-central regions

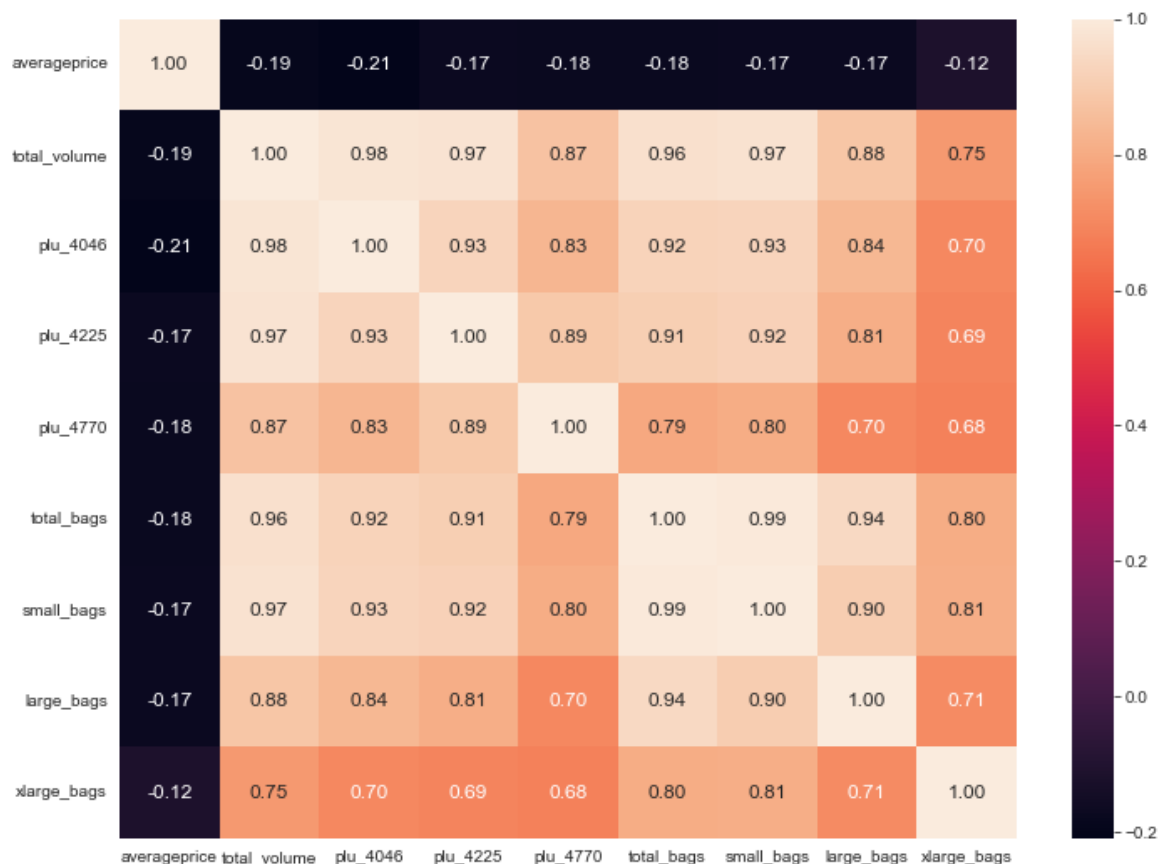
In 2017:

- average price of avocado was higher in hart fort Springfield, New York, San Francisco, charlotte, Sacramento regions. price was above 1.75 per unit.
- Houston, Dallas. South Central & Nashville had lower price per unit.
- Volume consumption was higher in California, West & South-central regions

In 2018:

- Data was only of 3 months, so it won't make sense to compare with other years data.

10. Correlation matrix of numerical features;



Observation:

- All Numerical features seem to have very low correlation with target variable 'Average Price' & high correlation with each other, hinting about multicollinearity.

Pre-Processing Data.

1. In this step we will make week number, month number & year as categorical datatypes.

Copying the dataframe & converting all the necessary features in to proper datatypes;

```
1 #Copying dataframe
2 data2 = data0.copy()
3
4 data2['date'] = pd.to_datetime(data2['date'], errors='raise')
5 data2['week_number'] = data2['date'].dt.week
6 data2['month_number'] = data2['date'].dt.month
7 data2['year'] = data2['date'].dt.year
8
9 data2.drop('date', axis=1, inplace=True)
10
11 data2['year'] = data2['year'].astype('str')
12 data2['week_number'] = data2['week_number'].astype('str')
13 data2['month_number'] = data2['month_number'].astype('str')
```

2. Checking Skewness & Removing using Power transformation

```
1 data2[['total_volume', 'plu_4046', 'plu_4225', 'plu_4770',
2       'total_bags', 'small_bags', 'large_bags']].skew()

total_volume    9.007687
plu_4046         8.648220
plu_4225         8.942466
plu_4770        10.159396
total_bags       9.756072
small_bags       9.540660
large_bags       9.796455
dtype: float64
```

Observation

1. All the numerical features have skewness. we will remove it using yeo-johnson transformation

```
1 x = data2[['total_volume', 'plu_4046', 'plu_4225', 'plu_4770',
2       'total_bags', 'small_bags', 'large_bags']]
```

```
1 from sklearn.preprocessing import PowerTransformer
2 pt = PowerTransformer(method='yeo-johnson', standardize=True)
3
4 x_pt = pt.fit_transform(x)
5 x_pt = pd.DataFrame(x_pt, columns = x.columns)
```

```
1 num_feat = ['total_volume', 'plu_4046', 'plu_4225', 'plu_4770', 'total_bags', 'small_bags', 'large_bags']
2 for col in num_feat:
3     data2[col]=x_pt[col]
```

```
1 data2[['total_volume', 'plu_4046', 'plu_4225', 'plu_4770',
2       'total_bags', 'small_bags', 'large_bags']].skew()

total_volume    0.011171
plu_4046        -0.022679
plu_4225        -0.017740
plu_4770         0.032184
total_bags      -0.001626
small_bags       0.007058
large_bags      -0.073809
dtype: float64
```

Observation


1. skewness in all numerical features is removed & are well within skewness of +/-0.5.

3. Checking outliers & removing if any using IQR Method

```
1 data3 = data2.copy()
```

```
1 data_out = data2[['total_volume', 'plu_4046', 'plu_4225', 'plu_4770',  
2 'total_bags', 'small_bags', 'large_bags']]
```

```
1 data_out.iloc[:,:].boxplot(figsize = (16,8))  
2 plt.subplots_adjust(bottom=0.25)  
3 plt.show()
```



```
25 # creating a function to remove outliers using inter quantile range...  
26 def get_outliers_iqr(data):  
27  
28     q1 = data.quantile(0.25)  
29     q3 = data.quantile(0.75)  
30     iqr = q3-q1  
31  
32     #empty list to store index values  
33     all_indices = []  
34  
35     for column in data.columns:  
36         lower, upper = (q1[column] - (1.5*iqr[column])), (q3[column] + (1.5*iqr[column]))  
37  
38         index = np.where((data[column] < lower) | (data [column] > upper))  
39  
40         all_indices.extend(index[0])  
41  
42     set_res = set(all_indices)  
43     indices_to_remove = np.array(list(set_res))  
44     return indices_to_remove
```

```
1 data_out = data2[['total_volume', 'plu_4046', 'plu_4225', 'plu_4770',  
2 'total_bags', 'small_bags', 'large_bags']]
```

```
1 to_remove_iqr = get_outliers_iqr(data_out)  
2 data_out = data_out.drop(data_out.index[to_remove_iqr])  
3 data3_iqr = data3.drop(data3.index[to_remove_iqr])
```

```
1 total_data_loss = (data3.shape[0] - data3_iqr.shape[0])*100/data3.shape[0]  
2 print('Total Data loss after removing Outliers : ', round(total_data_loss,2))
```

Total Data loss after removing Outliers : 1.9


Observation:

- We see that total data loss after removing outliers was 1.9%.

```

1 data_out.iloc[:,:].boxplot(figsize = (16,8))
2 plt.subplots_adjust(bottom=0.25)
3 plt.show()

```



Observation

1. using IQR Method, outliers removed from all the columns
2. Data loss after removing outliers using IQR is : 1.9%

4. We will Remove all the observations where region column has 'TotalUS' entry;

```

1 data5 = data3_iqr.copy()
2 data6 = data5.drop(columns = ['total_bags'])
3 data6 = data6.loc[data6['region'] != 'TotalUS']
4 data6

```

	averageprice	total_volume	plu_4046	plu_4225	plu_4770	small_bags	large_bags	xlarge_bags	type	year	region	week_number
0	1.33	-0.086605	-0.658619	0.307313	-0.247250	-0.340627	-0.773351	0.0	conventional	2015	Albany	52
1	1.35	-0.155746	-0.776870	0.231647	-0.201041	-0.306425	-0.763435	0.0	conventional	2015	Albany	51
2	0.93	0.179327	-0.731947	0.579503	-0.007072	-0.366303	-0.750820	0.0	conventional	2015	Albany	50
3	1.08	0.003876	-0.634200	0.415100	-0.148322	-0.496853	-0.691865	0.0	conventional	2015	Albany	49
4	1.28	-0.187640	-0.685293	0.224806	-0.137918	-0.477201	-0.600975	0.0	conventional	2015	Albany	48
...
18244	1.63	-0.674587	-0.467366	-0.919470	-1.238310	-0.178839	-0.410859	0.0	organic	2018	WestTexNewMexico	5
18245	1.71	-0.767550	-0.619883	-0.665372	-1.238310	-0.325977	-0.481465	0.0	organic	2018	WestTexNewMexico	4
18246	1.87	-0.771505	-0.619831	-0.772496	0.401546	-0.308728	-0.943486	0.0	organic	2018	WestTexNewMexico	3
18247	1.93	-0.698077	-0.550282	-0.710543	0.401246	-0.248932	-0.908510	0.0	organic	2018	WestTexNewMexico	2
18248	1.62	-0.663809	-0.367908	-0.785173	0.122919	-0.212590	-1.042250	0.0	organic	2018	WestTexNewMexico	1

17725 rows × 13 columns

5. Using Label encoder for encoding region column.

Sklearn provides a very efficient tool for encoding the levels of categorical features into numeric values. Label Encoder encode labels with a value between 0 and n_classes-1 where n is the number of distinct labels.


```
1 from sklearn.preprocessing import LabelEncoder
2 le = LabelEncoder()
3 data6['region']=le.fit_transform(data6['region'])
4 data7 = pd.get_dummies(data6, drop_first = True)
5 data7
```

	averageprice	total_volume	plu_4046	plu_4225	plu_4770	small_bags	large_bags	xlarge_bags	region	type_organic	year_2016	year_2017	year_2018
0	1.33	-0.086605	-0.658619	0.307313	-0.247250	-0.340627	-0.773351	0.0	0	0	0	0	0
1	1.35	-0.155746	-0.776870	0.231647	-0.201041	-0.306425	-0.763435	0.0	0	0	0	0	0
2	0.93	0.179327	-0.731947	0.579503	-0.007072	-0.366303	-0.750820	0.0	0	0	0	0	0
3	1.08	0.003876	-0.634200	0.415100	-0.148322	-0.496853	-0.691865	0.0	0	0	0	0	0
4	1.28	-0.187640	-0.685293	0.224806	-0.137918	-0.477201	-0.600975	0.0	0	0	0	0	0
...
18244	1.63	-0.674587	-0.467366	-0.919470	-1.238310	-0.178839	-0.410859	0.0	52	1	0	0	1
18245	1.71	-0.767550	-0.619883	-0.665372	-1.238310	-0.325977	-0.481465	0.0	52	1	0	0	1
18246	1.87	-0.771505	-0.619831	-0.772496	0.401546	-0.308728	-0.943486	0.0	52	1	0	0	1
18247	1.93	-0.698077	-0.550282	-0.710543	0.401246	-0.248932	-0.908510	0.0	52	1	0	0	1
18248	1.62	-0.663809	-0.367908	-0.785173	0.122919	-0.212590	-1.042250	0.0	52	1	0	0	1

17725 rows x 76 columns

6. We will check for correlation & Multicollinearity using .corr() method & VIF resply.

```
1 x_check = data7[['total_volume', 'plu_4046', 'plu_4225', 'plu_4770', 'small_bags', 'large_bags', 'xlarge_bags', 'averageprice']]
2 x_check.corr()
3 plt.figure(figsize=(8,5))
4 sns.heatmap(x_check.corr(),annot=True, fmt = "%.2f", annot_kws={'size':12});
5 plt.show()
```



Observation:

- We see there are some features with highly correlated with each other, We will confirm the same using Statsmodels' OLS result as follows.

```

1 model = sm.OLS(x_check['averageprice'], sm.add_constant(x_check.drop('averageprice', axis=1))).fit()
2 print(model.summary())

```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          averageprice    R-squared:                0.438
Model:                  OLS            Adj. R-squared:           0.438
Method:                 Least Squares   F-statistic:              1973.
Date:                   Wed, 05 Oct 2022 Prob (F-statistic):       0.00
Time:                   23:29:06        Log-Likelihood:          -3935.4
No. Observations:       17725          AIC:                   7887.
Df Residuals:           17717          BIC:                   7949.
Df Model:                7
Covariance Type:        nonrobust
=====
                    coef    std err          t      P>|t|      [0.025     0.975]
-----
const                1.3947     0.002    592.150     0.000     1.390     1.399
total_volume        -0.6076     0.014   -44.473     0.000    -0.634    -0.581
plu_4046             0.0260     0.006     4.656     0.000     0.015     0.037
plu_4225             0.3082     0.007    42.910     0.000     0.294     0.322
plu_4770            -0.0721     0.004   -17.143     0.000    -0.080    -0.064
small_bags           0.1650     0.007    23.515     0.000     0.151     0.179
large_bags          -0.0506     0.003   -15.377     0.000    -0.057    -0.044
xlarge_bags         1.907e-06   3.03e-07     6.303     0.000   1.31e-06   2.5e-06
=====
Omnibus:                 1712.926   Durbin-Watson:           0.253
Prob(Omnibus):            0.000   Jarque-Bera (JB):        2638.681
Skew:                     0.728   Prob(JB):                 0.00
Kurtosis:                  4.206   Cond. No.                  6.00e+04
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 6e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Observation:

- Condition number is more than 10, this indicate there is multicollinearity issue.

7. We will remove multicollinearity using Variance inflation factor

```

1 from statsmodels.stats.outliers_influence import variance_inflation_factor
2 df_feat = x_check.drop('averageprice', axis=1)
3 df_tgt = x_check['averageprice']
4
5 vif = pd.DataFrame()
6 vif["variables"] = df_feat.columns
7 vif["VIF"] = [variance_inflation_factor(df_feat.values, i) for i in range(df_feat.shape[1])]
8 vif.sort_values(by = 'VIF', ascending = False)

```

	variables	VIF
0	total_volume	34.198010
2	plu_4225	9.428277
4	small_bags	8.120107
1	plu_4046	5.774840
3	plu_4770	3.349894
5	large_bags	1.986215
6	xlarge_bags	1.184613

Observation:

- We will remove total volume as it shows VIF more than 10.

```

1 df_feat = x_check.drop(['averageprice', 'total_volume'], axis=1)
2 df_tgt = x_check['averageprice']
3
4 vif = pd.DataFrame()
5 vif["variables"] = df_feat.columns
6 vif["VIF"] = [variance_inflation_factor(df_feat.values, i) for i in range(df_feat.shape[1])]
7 vif.sort_values(by = 'VIF', ascending = False)

```

	variables	VIF
1	plu_4225	4.020856
3	small_bags	3.862361
0	plu_4046	3.432350
2	plu_4770	3.349892
4	large_bags	1.763164
5	xlarge_bags	1.164698

Observation

1. Multicollinearity is removed.. VIF for all the features is within 10.

8. Final data preparation for model training.

We will separate dependant & independent features.

Preparing data for ML model

```

1 x = data7.drop(columns = ['region', 'averageprice', 'total_volume'])
2 y_class = data7['region']
3 y_reg = data7['averageprice']

```

```
1 x
```

	plu_4046	plu_4225	plu_4770	small_bags	large_bags	xlarge_bags	type_organic	year_2016	year_2017	year_2018	week_number_10	week_number_11
0	-0.658619	0.307313	-0.247250	-0.340627	-0.773351	0.0	0	0	0	0	0	(
1	-0.776870	0.231647	-0.201041	-0.306425	-0.763435	0.0	0	0	0	0	0	(
2	-0.731947	0.579503	-0.007072	-0.366303	-0.750820	0.0	0	0	0	0	0	(
3	-0.634200	0.415100	-0.148322	-0.496853	-0.691865	0.0	0	0	0	0	0	(
4	-0.685293	0.224806	-0.137918	-0.477201	-0.600975	0.0	0	0	0	0	0	(
...
18244	-0.467366	-0.919470	-1.238310	-0.178839	-0.410859	0.0	1	0	0	1	0	(
18245	-0.619883	-0.665372	-1.238310	-0.325977	-0.481465	0.0	1	0	0	1	0	(
18246	-0.619831	-0.772496	0.401546	-0.308728	-0.943486	0.0	1	0	0	1	0	(
18247	-0.550282	-0.710543	0.401246	-0.248932	-0.908510	0.0	1	0	0	1	0	(
18248	-0.367908	-0.785173	0.122919	-0.212590	-1.042250	0.0	1	0	0	1	0	(


17725 rows × 73 columns

9. Checking data imbalance for classification problem

```

1 sns.countplot(x=y_class)
2 plt.show()

```



No need to do use SMOTE method as the data is perfectly balanced.

10. Applying Standard Scaling to Independent features

using Standard Scaler to standardize the data

```
1 #Transforming data into Standard Normal Distribution
2 from sklearn.preprocessing import StandardScaler
3 sc = StandardScaler()
```

```
1 X = pd.DataFrame(sc.fit_transform(x), columns=x.columns)
2 X
```

	plu_4046	plu_4225	plu_4770	small_bags	large_bags	xlarge_bags	type_organic	year_2016	year_2017	year_2018	week_number_10	week_number_11
0	-0.647482	0.347429	-0.234376	-0.358974	-0.764661	-0.241634	-0.989393	-0.668535	-0.680842	-0.278049	-0.1556	-0.155409
1	-0.768411	0.269506	-0.187556	-0.321898	-0.754458	-0.241634	-0.989393	-0.668535	-0.680842	-0.278049	-0.1556	-0.155409
2	-0.722471	0.627737	0.008975	-0.386808	-0.741477	-0.241634	-0.989393	-0.668535	-0.680842	-0.278049	-0.1556	-0.155409
3	-0.622510	0.458430	-0.134141	-0.528330	-0.680813	-0.241634	-0.989393	-0.668535	-0.680842	-0.278049	-0.1556	-0.155409
4	-0.674761	0.262461	-0.123600	-0.507026	-0.587286	-0.241634	-0.989393	-0.668535	-0.680842	-0.278049	-0.1556	-0.155409
...
17720	-0.451898	-0.915943	-1.238529	-0.183588	-0.391655	-0.241634	1.010720	-0.668535	-0.680842	3.596491	-0.1556	-0.155409
17721	-0.607869	-0.654267	-1.238529	-0.343093	-0.464309	-0.241634	1.010720	-0.668535	-0.680842	3.596491	-0.1556	-0.155409
17722	-0.607816	-0.764586	0.422992	-0.324393	-0.939732	-0.241634	1.010720	-0.668535	-0.680842	3.596491	-0.1556	-0.155409
17723	-0.536692	-0.700785	0.422687	-0.259572	-0.903741	-0.241634	1.010720	-0.668535	-0.680842	3.596491	-0.1556	-0.155409
17724	-0.350187	-0.777641	0.140683	-0.220177	-1.041360	-0.241634	1.010720	-0.668535	-0.680842	3.596491	-0.1556	-0.155409

17725 rows × 13 columns

Building Machine Learning Models.

Please refer the notebook in the [github link](#) for detailed codes:

For Classification model;

We will test the following models (the detailed code has all algorithm run in a loop to get best random state & best cross validation fold). The code will generate a table consisting best random state & Cross validation fold & all related evaluation metrics of the respective model.

```
1 dtc = DecisionTreeClassifier()
2 etc = ExtraTreesClassifier()
3 knn = KNeighborsClassifier()
4 lgr = LogisticRegression(multi_class='ovr')
5 rfc = RandomForestClassifier()

# To print classification report, confusion matrix, roc-auc curve
def print_score(clf, x_train, x_test, y_train, y_test, train=True):
    if train:
        y_pred = clf.predict(x_train)
        print(f"Accuracy score for Train : {accuracy_score(y_train, y_pred) * 100:.2f}%")
        return round(accuracy_score(y_train, y_pred) * 100, 2)
    #
    print('\n \n Train Classification report \n', classification_report(y_train, y_pred, digits=2))

    elif train==False:
        y_pred = clf.predict(x_test)
        acc = round(accuracy_score(y_test, y_pred) * 100, 2)
        print(f"Accuracy score for Test : {accuracy_score(y_test, y_pred) * 100:.2f}%")
        confusion_matrix_c(y_test, y_pred)
        print('\n \n Test Classification report \n', classification_report(y_test, y_pred, digits=2))

        diff = []
        best_cv = []
        for j in range(3, 12):
            cv_score = round(cross_val_score(clf, X, y_class, cv=j, scoring="accuracy").mean() * 100, 2)
            diff.append(abs(cv_score - acc))
            best_cv.append(j)
        k_f = best_cv[diff.index(min(diff))]
        cv_score = cross_val_score(clf, X, y_class, cv=k_f, scoring="accuracy").mean()
        print(f"Cross Validation score at best cv={k_f} is : {cv_score * 100:.2f}%")
        return acc, k_f, min(diff), cv_score * 100

models = [dtc, etc, knn, lgr, rfc]
models_name = ['DecisionTreeClassifier', 'ExtraTreesClassifier', 'KNeighborsClassifier',
               'LogisticRegression', 'RandomForestClassifier']

dummy_count = 0 #dummy variable for count purpose
for model in models:
    ### splitting with best random state
    x_train, x_test, y_train, y_test = train_test_split(X, y_class, random_state=best_i, test_size=.2, stratify=y_class)

    ### training the model
    model.fit(x_train, y_train)
    algo.append(models_name[dummy_count])
    tr_ac = print_score(model, x_train, x_test, y_train, y_test, train=True)
```

	algo	best random state	train_accuracy	test_accuracy	Score_diff	best cv fold	cross_val_score
3	LogisticRegression	78	40.97	30.69	0.20	9	30.894866
2	KNeighborsClassifier	60	22.03	0.20	9.54	3	9.737851
0	DecisionTreeClassifier	49	100.00	79.10	15.10	11	63.773091
4	RandomForestClassifier	82	100.00	85.42	15.72	11	69.453270
1	ExtraTreesClassifier	123	100.00	77.04	15.91	11	60.838086

Observation:

- we will select Random Forest classifier as it has higher test accuracy & lower difference between test accuracy & Cross value score amongst all models.
- It also has higher test accuracy of all the models tested.

Hyper parameter tuning for classification model;

```
1 param_grid_rfr = {'n_estimators':[200,300,350,450],
2                  'max_depth':[50, 65, 75, 80],
3                  'min_samples_split':[2, 3, 5],
4                  'criterion':['gini', 'entropy']}
5
```

```
1 rfc_tune = RandomForestClassifier()
```

```
1 x_train, x_test, y_train, y_test = train_test_split(X, y_class, test_size = 0.2, random_state = 82, stratify=y_class)
```

```
1 rf_grid = GridSearchCV(estimator = rfc_tune,
2                        param_grid = param_grid_rfr,
3                        verbose = 2,
4                        scoring = 'accuracy')
```

```
1 rf_grid.fit(x_train, y_train)
```

...

```
1 rf_grid.best_params_
```

```
{'criterion': 'entropy',
 'max_depth': 80,
 'min_samples_split': 3,
 'n_estimators': 450}
```

```
1 rf_grid.best_score_
```

```
0.8555007052186177
```

Using hyper parameters obtained from gridsearch for classificaion model...

```
1 x_train, x_test_class, y_train, y_test_class = train_test_split(X, y_class, test_size = 0.2,
2                                                                random_state = 82, stratify=y_class)
```

```
1 rfc_tune_final = RandomForestClassifier(criterion='entropy', max_depth= 80, min_samples_split= 3, n_estimators= 450)
```

```
1 rfc_tune_final.fit(x_train, y_train)
2 y_pred=rfc_tune_final.predict(x_test_class)
3 print('Accuracy Score: ', accuracy_score(y_test_class, y_pred))
```

```
Accuracy Score: 0.8750352609308886
```

```

1 confusion_matrix_c(y_test_class, y_pred)
2 print('\n \n Test Classification report \n', classification_report(y_test_class, y_pred, digits=2))
3
4 cv_score = cross_val_score(rfc_tune_final,X, y_class, cv=11, scoring="accuracy").mean()
5 print(f"Cross Validation score at best cv = 11 is : {cv_score*100:.2f}%")

```

=====
Confusion Matrix :

```

[[63  0  0 ...  0  0  0]
 [ 0 50  0 ...  0  0  0]
 [ 0  0 65 ...  0  0  0]
 ...
 [ 0  0  0 ... 23  0  0]
 [ 0  0  0 ...  0 68  0]
 [ 0  0  0 ...  0  0 52]]
=====

```

```

      accuracy          0.88      3545
macro avg      0.88      0.87      0.87      3545
weighted avg   0.88      0.88      0.87      3545

```

Cross Validation score at best cv = 11 is : 73.17%

observation:

After tuning we see increase in test accuracy & cross value scores. Hence our model is good.

Saving & predicting classification model

```

1 filename='avocado_class.pkl'
2 pickle.dump(rfc_tune_final,open(filename,'wb'))
3
4 filename='avocado_reg.pkl'
5 pickle.dump(xgb_tune_final,open(filename,'wb'))

```

```

1 model =pickle.load(open('avocado_class.pkl','rb'))
2 pred =model.predict(x_test_class)
3 result = pd.DataFrame(list(zip(y_test_class, pred)), columns = ['Actual', 'Predicted'])
4 result

```

	Actual	Predicted
0	35	35
1	12	12
2	2	2
3	8	8
4	27	27
...
3540	27	27
3541	3	3
3542	33	33
3543	23	3
3544	47	47

3545 rows × 2 columns

For Regression model;

We will test the following models (the detailed code has all algorithm run in a loop to get best random state & best cross validation fold). The code will generate a table consisting best random state & Cross validation fold & all related evaluation metrics of the respective model.

```
1 lr = LinearRegression()
2 ls = Lasso()
3 rd = Ridge()
4 rfr = RandomForestRegressor()
5 abr = AdaBoostRegressor()
6 gbr = GradientBoostingRegressor()
7 dtr = DecisionTreeRegressor()
8 knr = KNeighborsRegressor()
9 xgb = XGBRegressor()

models = [lr, ls, rd, abr, gbr, dtr, knr, xgb, rfr]
models_name = ['Linear Regression', 'Lasso', 'Ridge',
               'Ada-Boost Regressor', 'Gradient Boosting Regressor',
               'Decision Tree Regressor', 'KNeighbors Regressor',
               'XGB Regressor', 'Random Forest Regressor']
dummy_count = 0 #dummy variable for count purpose
for model in models:
    ### splitting the train7 set with best random state
    x_train, x_test, y_train, y_test = train_test_split(X, y_reg, random_state=best_i, test_size=.25)

    ### training the model
    model.fit(x_train, y_train)
    algo.append(models_name[dummy_count])

    ##showing the results in output
    print('::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::')
    print(' ')
    print(' ')
    print(models_name[dummy_count] + ' Model')
    print(' ')
    print('for ' + models_name[dummy_count] + ' model, Best Random_state number for splitting the data is: ', best_i)
    print(' ')
    print('===scores for training set===')
    print('r2 score for training set', r2_score(y_train, pred_train))
    print('MAE for training set: ', mean_absolute_error(y_train, pred_train))
    print('MSE for training set: ', mean_squared_error(y_train, pred_train))
    print('SMSE for training set: ', np.sqrt(mean_squared_error(y_train, pred_train)))
    print(' ')
    print('===scores for testing set===')
    print('r2 score for testing set : ', r2_score(y_test, pred_test))
    print('MAE for testing set: ', mean_absolute_error(y_test, pred_test))
    print('MSE for testing set: ', mean_squared_error(y_test, pred_test))
    print('SMSE for testing set: ', np.sqrt(mean_squared_error(y_test, pred_test)))
    print(' ')
    print(' ')
```

```

### printing CV Score based on best CV fold value
k_f = best_cv[cv_diff.index(min(cv_diff))]
cv_fold.append(k_f)
cv_score = cross_val_score(model, X, y_reg, cv=k_f).mean()
crossvalscore.append(round(cv_score,2))
score_diff.append(abs(round(cv_score - test_accuracy,2)))
print(f"Cross Validation score at best cv={k_f} is : {cv_score*100:.2f}%")
dummy_count+=1
print(' ')
print(' ')

```

Sr. No.	algo	best random state	train_r2score	test_r2score	train_mae	train_mse	train_rmse	test_mae	test_mse	test_rmse	best cv fold	cross_val_score
9	Random Forest Regressor	92	0.98	0.86	0.04	0.00	0.06	0.11	0.02	0.15	11	0.39
8	XGB Regressor	63	0.93	0.85	0.08	0.01	0.10	0.12	0.03	0.16	12	0.51
6	Decision Tree Regressor	18	1.00	0.71	0.00	0.00	0.00	0.14	0.05	0.22	11	-0.17
5	Gradient Boosting Regressor	76	0.71	0.71	0.16	0.05	0.22	0.16	0.05	0.21	9	0.38
7	KNeighbors Regressor	1	0.76	0.65	0.15	0.04	0.20	0.18	0.06	0.25	12	0.33
1	Linear Regression	20	0.57	0.57	0.21	0.07	0.27	0.20	0.07	0.26	9	0.15
3	Ridge	20	0.57	0.57	0.21	0.07	0.27	0.20	0.07	0.26	9	0.15
4	Ada-Boost Regressor	9	0.52	0.51	0.22	0.08	0.28	0.23	0.08	0.28	10	-0.16
2	Lasso	23	0.00	-0.00	0.32	0.16	0.40	0.33	0.16	0.40	9	-1.18

Observation:

- we will select XGB Regressor as it has higher test & Cross validation scores, also it has lower MAE value.
- Random forest has higher test score but cross value score is very less. Hence, we will use XGB Regressor.

Hyper parameter tuning for regression model;

Using hyper parameters obtained from gridsearch for regression model...

```
1 x_train, x_test_reg, y_train, y_test_reg = train_test_split(X, y_reg, random_state=63, test_size=.25)
```

```
1 xgb_tune_final = XGBRegressor(learning_rate= 0.1, max_depth=8, n_estimators=800)
```

```
1 xgb_tune_final.fit(x_train,y_train)
2 y_pred=xgb_tune_final.predict(x_test_reg)
```


```
1 print('r2 score for testing set : ', r2_score(y_test_reg, y_pred))
2 print('MAE for testing set: ', mean_absolute_error(y_test_reg, y_pred))
3 print('MSE for testing set: ', mean_squared_error(y_test_reg, y_pred))
4 print('SMSE for testing set: ', np.sqrt(mean_squared_error(y_test_reg, y_pred)))
```

r2 score for testing set : 0.874351595012883
 MAE for testing set: 0.10166876581386539
 MSE for testing set: 0.02122898504075049
 SMSE for testing set: 0.14570169882589046

```

1 ##plotting the graph with bestfit Line, actual & predicted values
2 plt.figure(figsize = (6,5))
3 plt.scatter(x = y_test_reg, y=y_pred, color = 'r')
4 plt.plot(y_test_reg, y_test_reg, color = 'b')
5 plt.xlabel('actual', fontsize = 15)
6 plt.ylabel('predictions', fontsize = 15)
7 plt.title('Hyper parameter tuned Lasso Regressor', fontsize = 15)
8 plt.show()

```



```
1 cross_val_score(xgb_tune_final, X, y_reg, cv = 12, scoring = 'r2').mean()
```

0.5139095135730846

observation:

After tuning we see increase in test score & cross value scores. Hence our model is good.

Saving & predicting regression model

```
1 model = pickle.load(open('avocado_reg.pkl', 'rb'))
2 pred = model.predict(x_test_reg)
3 result = pd.DataFrame(list(zip(y_test_reg, pred)), columns = ['Actual', 'Predicted'])
4 result
```

	Actual	Predicted
0	1.60	1.655651
1	1.05	1.003266
2	1.78	1.507349
3	1.61	1.596217
4	1.60	1.697207
...
4427	1.26	1.178945
4428	1.26	1.179238
4429	1.55	1.754191
4430	1.17	1.011211
4431	1.70	1.568804

4432 rows × 2 columns